

Project2 Write Up

1(a)

i.

```
[ (445_p2) mason-wudeMacBook-Pro:probuzhidaogaigaishamingzi:project2 mason_wu$ python ]
dataset.py
loading train...
loading val...
loading test...
Train:  300
Val:    150
Test:   100
Mean: [123.968 117.887  93.942]
Std:  [63.033 59.667 61.949]
```

ii.

If statistics are extracted from a validation or test set, data leakage is introduced, causing the model to behave on the test set in a "preference" for the statistical distribution of that data set, which does not accurately reflect the generalization ability of the model. Therefore, using the mean and standard deviation of the training set helps to ensure that the validation set and the test set remain "independent," thus testing the model's performance on real, unseen data.

1(b)



2(a)

Max Pooling layers don't have parameters.

Layer 1: Convolutional Layer 1

Input channels: 3

Output channels: 16

Parameters per convolution kernel: $5 \times 5 = 25$

Number of weight parameters: $3 \times 16 \times 25 = 1200$

Number of bias parameters: 16

Total parameters: $1200+16=1216$

Layer 3: Convolutional Layer 2

Input channels: 16

Output channels: 64

Parameters per convolution kernel: $5 \times 5 = 25$

Number of weight parameters: $16 \times 64 \times 25 = 25600$

Number of bias parameters: 64

Total parameters: $25600+64=25664$

Layer 5: Convolutional Layer 3

Input channels: 64

Output channels: 8

Parameters per convolution kernel: $5 \times 5 = 25$

Number of weight parameters: $64 \times 8 \times 25 = 12800$

Number of bias parameters: 8

Total parameters: $12800+8=12808$

Layer 6: Fully Connected Layer 1 (Output Layer)

Input size: 32 (output from previous layer $8 \times 2 \times 2$, the two 2 here is because we reduce the image size by 2 through pooling twice and we need to resize it back)

Output size: 2 (two categories)

Number of weight parameters: $32 \times 2 = 64$

Number of bias parameters: 2

Total parameters: $64+2=66$

In total, we have $1216+25664+12808+66=39754$ parameters.

2(b)

2(c)

2(d)

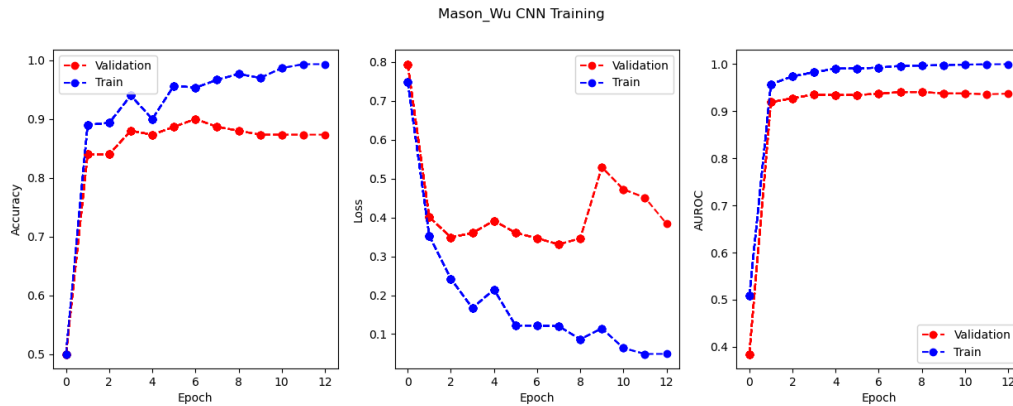
2(e)

2(f)

i.

```
(445_p2) mason-wudeMacBook-Pro:probuzhidaogaigaishamingzi:project2 mason
_wu$ python train_cnn.py
loading train...
loading val...
loading test...
Number of float-valued parameters: 39754
Loading cnn...
No saved model parameters found
Epoch 0
```

```
Epoch 0
  Validation Accuracy:0.5
  Validation Loss:0.7946
  Validation AUROC:0.3849
  Train Accuracy:0.5
  Train Loss:0.7496
  Train AUROC:0.5083
Epoch 1
  Validation Accuracy:0.84
  Validation Loss:0.4011
  Validation AUROC:0.9191
  Train Accuracy:0.89
  Train Loss:0.353
  Train AUROC:0.9572
Epoch 2
  Validation Accuracy:0.84
  Validation Loss:0.349
  Validation AUROC:0.9276
  Train Accuracy:0.8933
  Train Loss:0.2423
  Train AUROC:0.974
Epoch 3
  Validation Accuracy:0.88
  Validation Loss:0.36
  Validation AUROC:0.9355
  Train Accuracy:0.94
  Train Loss:0.1669
  Train AUROC:0.9827
Epoch 4
  Validation Accuracy:0.8733
  Validation Loss:0.3918
  Validation AUROC:0.9342
  Train Accuracy:0.9
  Train Loss:0.2147
  Train AUROC:0.9904
Epoch 5
  Validation Accuracy:0.8867
  Validation Loss:0.3607
  Validation AUROC:0.9346
  Train Accuracy:0.9567
  Train Loss:0.1222
  Train AUROC:0.9906
Epoch 6
  Validation Accuracy:0.9
  Validation Loss:0.3466
  Validation AUROC:0.9371
  Train Accuracy:0.9533
  Train Loss:0.1218
  Train AUROC:0.9926
Epoch 7
  Validation Accuracy:0.8867
  Validation Loss:0.3308
  Validation AUROC:0.9406
  Train Accuracy:0.9667
  Train Loss:0.1204
  Train AUROC:0.9963
Epoch 8
  Validation Accuracy:0.88
  Validation Loss:0.3462
  Validation AUROC:0.9408
  Train Accuracy:0.9767
  Train Loss:0.0866
  Train AUROC:0.997
Epoch 9
  Validation Accuracy:0.8733
  Validation Loss:0.5295
  Validation AUROC:0.9378
  Train Accuracy:0.97
  Train Loss:0.1147
  Train AUROC:0.998
Epoch 10
  Validation Accuracy:0.8733
  Validation Loss:0.4725
  Validation AUROC:0.9377
  Train Accuracy:0.9867
  Train Loss:0.0641
  Train AUROC:0.9988
Epoch 11
  Validation Accuracy:0.8733
  Validation Loss:0.4511
  Validation AUROC:0.936
  Train Accuracy:0.9933
  Train Loss:0.0487
  Train AUROC:0.9996
Epoch 12
  Validation Accuracy:0.8733
  Validation Loss:0.3841
  Validation AUROC:0.9372
  Train Accuracy:0.9933
  Train Loss:0.0489
  Train AUROC:0.9998
Finished Training
```



- Mini-batch Variability

Within each epoch, the training process splits the data set into multiple mini-batches for training. Each small batch of data does not fully represent the entire data distribution, which leads to a certain randomness in the update of model parameters. Therefore, the validation loss obtained by the model after each epoch may fluctuate, showing that the validation loss is not always monotonically decreasing. The randomness of small batches introduces inconsistent data samples, making the model parameters slightly different after each epoch, and this randomness causes validation losses to occasionally rise.

- Learning-rate Adjustment

In the training process, the model will use the set learning rate to update parameters, and a higher learning rate will make the model's parameter update amplitude larger, which will cause the loss to fluctuate in the process of decline (especially in the initial training stage). When the learning rate is high, the model may produce overshooting during optimization, resulting in a certain degree of fluctuation in validation loss.

It stops at epoch 12.

ii.

To be fair, I also start from epoch 0 in this case.


```
[(445_p2) mason-wudeMacBook-Pro:probuzhidaogaigaishamingzi:project2 mason]
_wu$ python train_cnn.py
loading train...
loading val...
loading test...
Number of float-valued parameters: 39754
Loading cnn...
Which epoch to load from? Choose in range [0, 12]. Enter 0 to train
from scratch.
>> 0
Checkpoint not loaded
Checkpoint successfully removed
Epoch 0
    Validation Accuracy:0.5
    Validation Loss:0.7946
    Validation AUROC:0.3849
    Train Accuracy:0.5
    Train Loss:0.7496
    Train AUROC:0.5083
Epoch 1
    Validation Accuracy:0.84
    Validation Loss:0.4011
    Validation AUROC:0.9191
    Train Accuracy:0.89
    Train Loss:0.353
    Train AUROC:0.9572
Epoch 2
    Validation Accuracy:0.84
    Validation Loss:0.349
    Validation AUROC:0.9276
    Train Accuracy:0.8933
    Train Loss:0.2423
    Train AUROC:0.974
Epoch 3
    Validation Accuracy:0.88
    Validation Loss:0.36
    Validation AUROC:0.9355
    Train Accuracy:0.94
    Train Loss:0.1669
    Train AUROC:0.9827
Epoch 4
    Validation Accuracy:0.8733
    Validation Loss:0.3918
    Validation AUROC:0.9342
    Train Accuracy:0.9
    Train Loss:0.2147
    Train AUROC:0.9904
Epoch 5
    Validation Accuracy:0.8867
    Validation Loss:0.3607
    Validation AUROC:0.9346
    Train Accuracy:0.9567
    Train Loss:0.1222
    Train AUROC:0.9906
Epoch 6
    Validation Accuracy:0.9
    Validation Loss:0.3466
    Validation AUROC:0.9371
    Train Accuracy:0.9533
    Train Loss:0.1218
    Train AUROC:0.9926
Epoch 7
    Validation Accuracy:0.8867
    Validation Loss:0.3308
    Validation AUROC:0.9406
    Train Accuracy:0.9667
    Train Loss:0.1204
```

```
Train AUROC:0.9963
Epoch 8
Validation Accuracy:0.88
Validation Loss:0.3462
Validation AUROC:0.9408
Train Accuracy:0.9767
Train Loss:0.0866
Train AUROC:0.997
Epoch 9
Validation Accuracy:0.8733
Validation Loss:0.5295
Validation AUROC:0.9378
Train Accuracy:0.97
Train Loss:0.1147
Train AUROC:0.998
Epoch 10
Validation Accuracy:0.8733
Validation Loss:0.4725
Validation AUROC:0.9377
Train Accuracy:0.9867
Train Loss:0.0641
Train AUROC:0.9988
Epoch 11
Validation Accuracy:0.8733
Validation Loss:0.4511
Validation AUROC:0.936
Train Accuracy:0.9933
Train Loss:0.0487
Train AUROC:0.9996
Epoch 12
Validation Accuracy:0.8733
Validation Loss:0.3841
Validation AUROC:0.9372
Train Accuracy:0.9933
Train Loss:0.0489
Train AUROC:0.9998
Epoch 13
Validation Accuracy:0.8867
Validation Loss:0.4092
Validation AUROC:0.9328
Train Accuracy:0.98
Train Loss:0.0557
Train AUROC:0.9996
Epoch 14
Validation Accuracy:0.8667
Validation Loss:0.5644
Validation AUROC:0.9271
Train Accuracy:0.98
Train Loss:0.0588
Train AUROC:1.0
Epoch 15
Validation Accuracy:0.86
Validation Loss:0.4295
Validation AUROC:0.9316
Train Accuracy:1.0
Train Loss:0.0371
Train AUROC:1.0
Epoch 16
Validation Accuracy:0.88
Validation Loss:0.5049
Validation AUROC:0.9344
```

```

Train Accuracy:0.9967
Train Loss:0.0187
Train AUROC:1.0
Epoch 17
Validation Accuracy:0.88
Validation Loss:0.4557
Validation AUROC:0.9364
Train Accuracy:1.0
Train Loss:0.0119
Train AUROC:1.0
Finished Training

```



It stops at epoch 17.

A patience value around 5 would be good: The loss curve of the verification set increases slightly at some epochs, but generally shows a downward trend. Therefore, if we choose an excessively small patience value (e.g. 2-3), we may cause the model to stop training prematurely, missing the opportunity for subsequent verification that the loss drops again. The fluctuation frequency of the verified loss curve in the chart is not high. A good patience number can be 5 or 6. This allows the model to continue training during a brief uptick in validation losses to see if subsequent validation losses drop again. If the patience value is too low, the opportunity to verify the overall decrease in set losses may be missed.

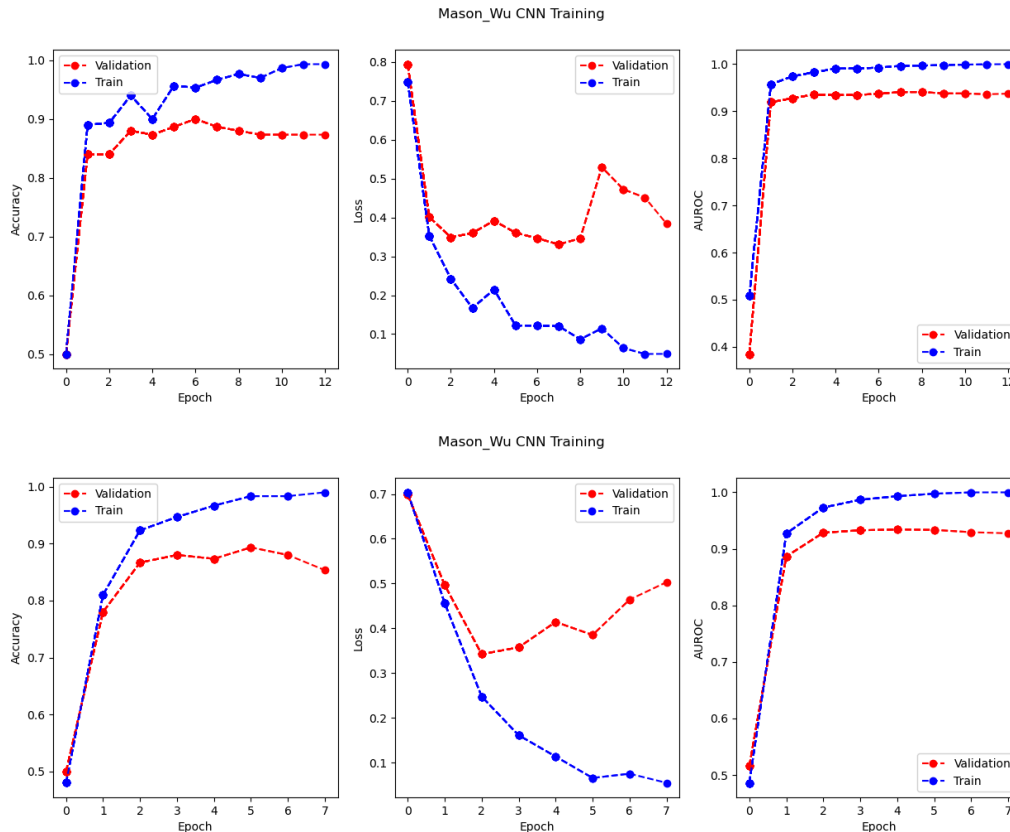
A larger patience value is appropriate for the following:

High noise data: If the fluctuations in validation losses are particularly frequent, a large patience value can prevent the model from stopping prematurely due to short-term fluctuations.

Long time training: When we have enough computing resources and time, we can increase patience and give the model longer time to find the optimal solution, thereby improving the final performance of the model.

iii.

The new size of the input to the fully connected layer is $64 \times 2 \times 2 = 256$.



	Epoch	Training AUROC	Validation AUROC
8 filters	7	0.9963	0.9406
64 filters	2	0.974	0.9276

If you're asking about the variance of the performance in this table:

As the number of filters in the convolutional layer increases from 8 to 64, we can observe the following performance changes:

1. Training AUROC: When the number of filters is increased to 64, the training AUROC is still very high (0.9975), close to 1, indicating that the model has a strong ability to distinguish on the training set.

2. Validation AUROC: Validation AUROC also decreased slightly, from 0.936 for 8 filters to 0.9335 for 64 filters. This shows that with the increase of the number of filters, the generalization ability of the model on the verification set does not improve significantly, and even slightly decreases.

One possible cause is that as the number of filters increases, so does the number of parameters in the model. This allows the model to fit more detail on the training set, but also increases the risk of overfitting. Therefore, although the training AUROC approaches 1, the validation AUROC does not increase or even decrease slightly.

If you're asking about the variance of the performance the whole graph:

- Training set performance:

8 filters: Training loss drops faster and more significantly at first and declines gradually later with a few rise-again. The accuracy and AUROC curve of the training set tend to be stable, the AUROC is around 1.

64 filters: Training loss drops less significantly with no rise-again, and the training set's AUROC is closer to 1. This shows that the model fits the training data better.

- Validation set performance:

8 filters: The validation loss leveled off after declining over the first few (about 8) epochs, but the validation AUROC hovered around about 0.9 without a significant improvement. The accuracy and AUROC of the validation set were stable.

64 filters: The validation loss decreased less, but fluctuated slightly in the middle and late stages, showing some instability. The maximum value of the validation set AUROC is also closer to 0.93.

Change reason analysis:

Model capacity increases: By increasing the number of filters from 8 to 64, the capacity of the model becomes larger and can capture more complex features, resulting in training loss and a slight decrease in AUROC. More filters make the model perform better on the training set, but may also lead to the risk of overfitting.

More stable performance of the validation set: The loss of the model with 64 filters on the validation set fluctuates less, indicating that as the model fits more features, these features generalize to the validation set, leading to a good fit on it.

Overall, increasing the number of filters improved the performance of the training set, but the performance of the validation set was not significantly improved.

```
[(445_p2) mason-wudeMacBook-Pro:probuzhidaogaigaishamingzi:project2 mason]
_wu$ python train_cnn.py
loading train...
loading val...
loading test...
Number of float-valued parameters: 129858
Loading cnn...
Which epoch to load from? Choose in range [0, 17]. Enter 0 to train
from scratch.
>> 0
Checkpoint not loaded
Checkpoint successfully removed
Epoch 0
    Validation Accuracy:0.5
    Validation Loss:0.6985
    Validation AUROC:0.5161
    Train Accuracy:0.48
    Train Loss:0.7038
    Train AUROC:0.4862
Epoch 1
    Validation Accuracy:0.78
    Validation Loss:0.4969
    Validation AUROC:0.8866
    Train Accuracy:0.81
    Train Loss:0.4564
    Train AUROC:0.9272
Epoch 2
    Validation Accuracy:0.8667
    Validation Loss:0.3424
    Validation AUROC:0.9285
    Train Accuracy:0.9233
    Train Loss:0.2478
    Train AUROC:0.9731
Epoch 3
    Validation Accuracy:0.88
    Validation Loss:0.3576
    Validation AUROC:0.9328
    Train Accuracy:0.9467
    Train Loss:0.1615
    Train AUROC:0.9868
Epoch 4
    Validation Accuracy:0.8733
    Validation Loss:0.4143
    Validation AUROC:0.9342
    Train Accuracy:0.9667
    Train Loss:0.1139
    Train AUROC:0.9931
Epoch 5
    Validation Accuracy:0.8933
    Validation Loss:0.3853
    Validation AUROC:0.9335
    Train Accuracy:0.9833
    Train Loss:0.066
    Train AUROC:0.9975
Epoch 6
    Validation Accuracy:0.88
    Validation Loss:0.4641
    Validation AUROC:0.9294
    Train Accuracy:0.9833
    Train Loss:0.0755
    Train AUROC:0.9998
Epoch 7
    Validation Accuracy:0.8533
    Validation Loss:0.5032
    Validation AUROC:0.9275
    Train Accuracy:0.99
    Train Loss:0.0551
```

```
Train AUROC:0.9999
Finished Training
```

2(g)

```
[{445_p2} mason-wudeMacBook-Pro:probuzhidaogaigaishamingzi:project2 mason]
_wu$ python test_cnn.py
loading train...
loading val...
loading test...
Loading cnn...
Which epoch to load from? Choose in range [0, 12]. Enter 0 to train
from scratch.
>> 7
Loading from checkpoint ./checkpoints/target/epoch=7.checkpoint.pth.
tar?
=> Successfully restored checkpoint (trained for 7 epochs)
Epoch 7
    Validation Accuracy:0.8867
    Validation Loss:0.3308
    Validation AUROC:0.9406
    Train Accuracy:0.9667
    Train Loss:0.1044
    Train AUROC:0.9963
    Test Accuracy:0.65
    Test Loss:1.31
    Test AUROC:0.6836
```


i.

	Training	Validation	Testing
Accuracy	0.9667	0.8867	0.65
AUROC	0.9963	0.9406	0.6836

ii.

We have evidence of overfitting here. (The training accuracy is 8 percentage points higher than the validation accuracy, the training AUROC is about 5.57 percentage points higher than the validation AUROC, indicating a significant gap.)

iii.



As you can see from the images, the data for the Positive label seems to be mostly collies (in Carly's kennel house), while the data for the Negative label is golden retrievers (in Paco's kennel house). This suggests that the labeling of positive and negative samples in the dataset is actually based on different dog breeds (Collie vs. Golden Retriever) to be tagged. Therefore, the model may learn the characteristics of different breeds of dogs, rather than other general characteristics.

Verification set performance: The accuracy of verification set was 87.33%, and the AUROC was 0.936, which performed relatively well.

Test set performance: The accuracy of the test set dropped to 64% with an AUROC of 0.6728, significantly lower than the performance of the validation set.

The model performs well on the training set and the validation set, but significantly poorly on the test set.

The probable causes of it are:

Inconsistent data distribution: The data distribution of the test set may be different, containing more uncommon variants or other breeds, resulting in a model that is difficult to identify correctly and performance deteriorates.

Overfitting: The model may overfit the features of the training set and the validation set during training, resulting in poor performance on test sets with large distribution differences.

3(a)

3(b)

3(c)



Epoch 11 corresponds to the model with the lowest validation loss.

```
((445_p2) mason-wudeMacBook-Pro:probuzhidaogaigaishamingzi:project2 mason
_wu$ python train_source.py
loading train...
loading val...
loading test...
Number of float-valued parameters: 39952
Loading source...
Which epoch to load from? Choose in range [0, 2]. Enter 0 to train f
rom scratch.
>> 0
Checkpoint not loaded
Checkpoint successfully removed
Epoch 0
    Validation Accuracy:0.1325
    Validation Loss:2.1531
    Validation AUROC:0.5084
    Train Accuracy:0.144
    Train Loss:2.148
    Train AUROC:0.5143
Epoch 1
    Validation Accuracy:0.1667
    Validation Loss:2.0443
    Validation AUROC:0.5963
    Train Accuracy:0.194
    Train Loss:2.03
    Train AUROC:0.6248
Epoch 2
    Validation Accuracy:0.2275
    Validation Loss:1.9745
    Validation AUROC:0.6643
    Train Accuracy:0.2415
    Train Loss:1.947
    Train AUROC:0.6912
Epoch 3
    Validation Accuracy:0.2642
    Validation Loss:1.9265
    Validation AUROC:0.6834
    Train Accuracy:0.302
    Train Loss:1.8707
    Train AUROC:0.7203
Epoch 4
    Validation Accuracy:0.3008
    Validation Loss:1.883
    Validation AUROC:0.7042
    Train Accuracy:0.3195
    Train Loss:1.7997
    Train AUROC:0.7523
Epoch 5
    Validation Accuracy:0.3033
    Validation Loss:1.8497
    Validation AUROC:0.7134
    Train Accuracy:0.3585
    Train Loss:1.7382
    Train AUROC:0.7724
Epoch 6
    Validation Accuracy:0.3075
    Validation Loss:1.8414
    Validation AUROC:0.7235
    Train Accuracy:0.3505
    Train Loss:1.7102
    Train AUROC:0.7919
Epoch 7
    Validation Accuracy:0.2933
    Validation Loss:1.831
    Validation AUROC:0.718
    Train Accuracy:0.4085
    Train Loss:1.6403
```

Epoch 8
Validation Accuracy:0.3192
Validation Loss:1.8111
Validation AUROC:0.7224
Train Accuracy:0.4185
Train Loss:1.5771
Train AUROC:0.8141

Epoch 9
Validation Accuracy:0.3117
Validation Loss:1.8214
Validation AUROC:0.7287
Train Accuracy:0.43
Train Loss:1.5566
Train AUROC:0.8268

Epoch 10
Validation Accuracy:0.3125
Validation Loss:1.8133
Validation AUROC:0.73
Train Accuracy:0.4565
Train Loss:1.5082
Train AUROC:0.8386

Epoch 11
Validation Accuracy:0.3325
Validation Loss:1.7831
Validation AUROC:0.7373
Train Accuracy:0.483
Train Loss:1.4339
Train AUROC:0.854

Epoch 12
Validation Accuracy:0.3233
Validation Loss:1.8486
Validation AUROC:0.7306
Train Accuracy:0.473
Train Loss:1.4112
Train AUROC:0.8637

Epoch 13
Validation Accuracy:0.32
Validation Loss:1.872
Validation AUROC:0.725
Train Accuracy:0.5095
Train Loss:1.3652
Train AUROC:0.8723

Epoch 14
Validation Accuracy:0.3283
Validation Loss:1.8206
Validation AUROC:0.7346
Train Accuracy:0.5565
Train Loss:1.2577
Train AUROC:0.8936


Epoch 15
Validation Accuracy:0.3233
Validation Loss:1.8203
Validation AUROC:0.7342
Train Accuracy:0.5995
Train Loss:1.191
Train AUROC:0.9087

Epoch 16
Validation Accuracy:0.3092
Validation Loss:1.9325
Validation AUROC:0.7272
Train Accuracy:0.574
Train Loss:1.2004
Train AUROC:0.9114

Epoch 17
Validation Accuracy:0.3242
Validation Loss:1.9675
Validation AUROC:0.7234

```
Train Accuracy:0.574
Train Loss:1.2004
Train AUROC:0.9114
Epoch 17
Validation Accuracy:0.3242
Validation Loss:1.9675
Validation AUROC:0.7234
Train Accuracy:0.6155
Train Loss:1.1295
Train AUROC:0.9237
Epoch 18
Validation Accuracy:0.3258
Validation Loss:1.9133
Validation AUROC:0.7253
Train Accuracy:0.6675
Train Loss:1.0131
Train AUROC:0.9379
Epoch 19
Validation Accuracy:0.3075
Validation Loss:1.9432
Validation AUROC:0.722
Train Accuracy:0.702
Train Loss:0.9299
Train AUROC:0.9503
Epoch 20
Validation Accuracy:0.2975
Validation Loss:2.1283
Validation AUROC:0.7174
Train Accuracy:0.695
Train Loss:0.9516
Train AUROC:0.9505
Epoch 21
Validation Accuracy:0.3108
Validation Loss:2.1175
Validation AUROC:0.7209
Train Accuracy:0.7275
Train Loss:0.8737
Train AUROC:0.9602
Finished Training
```

3(d)



Dalmatian(4), Saint Bernard(2), and Samoyed(0) are the most accurate.

Siberian Husky(7), Chihuahua(6), and Great Dane(3) are the least accurate.

Appearance similarity: Certain dog breeds have great similarity in appearance (e.g., coat color, body type, hair type, etc.), and models may be difficult to distinguish. For breeds such as Chihuahua and Samoyed, models can be confused due to their small size and white hair.

Data Imbalance: If there is insufficient sample size for certain breeds in the training data set, the model may perform poorly in these categories. Chihuahua has a higher number of misclassifications, possibly due to a smaller sample of the breed or a less distinct feature.

Feature complexity: Some dog breeds may have unique features that make it easier for models to learn. For example, Dalmatian spots are very distinct visual features, and models may perform better at recognizing them.

3(e)


3(f)

	AUROC		
	TRAIN	VAL	TEST
Freeze all CONV layers (Fine-tune FC layer)	0.8993	0.896	0.8204
Freeze first two CONV layers (Fine-tune last CONV and FC layers)	0.986	0.9081	0.768

Freeze first CONV layer (Fine-tune last 2 conv. and fc layers)	0.9964	0.9232	0.7196
Freeze no layers (Fine-tune all layers)	0.9872	0.9204	0.7292
No Pretraining or Transfer Learning (Section 2 performance)	0.9963	0.9406	0.6836

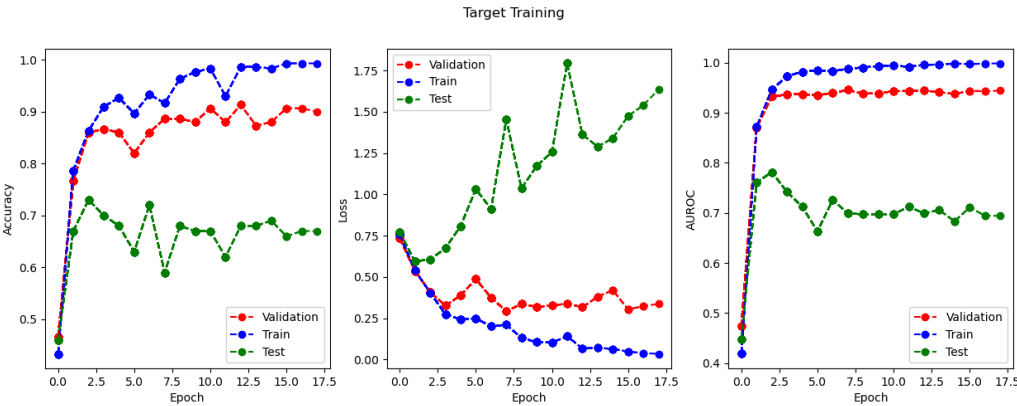
Epoch 2

Validation Accuracy:0.8667
 Validation Loss:0.3447
 Validation AUROC:0.9244
 Train Accuracy:0.91
 Train Loss:0.2608
 Train AUROC:0.9724
 Test Accuracy:0.69
 Test Loss:0.7385
 Test AUROC:0.7476



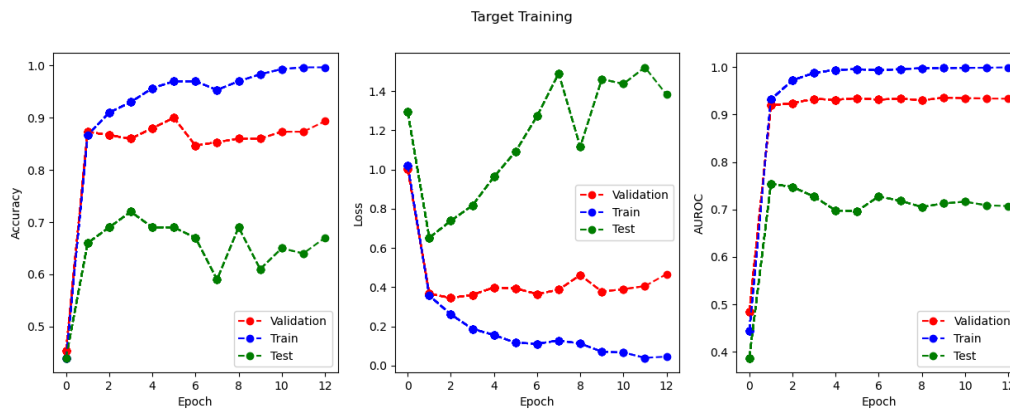
Epoch 3

Validation Accuracy:0.8667
 Validation Loss:0.3267
 Validation AUROC:0.9364
 Train Accuracy:0.91
 Train Loss:0.2744
 Train AUROC:0.9732
 Test Accuracy:0.7
 Test Loss:0.673
 Test AUROC:0.742



Epoch 2

Validation Accuracy:0.8667
Validation Loss:0.3447
Validation AUROC:0.9244
Train Accuracy:0.91
Train Loss:0.2608
Train AUROC:0.9724
Test Accuracy:0.69
Test Loss:0.7385
Test AUROC:0.7476



4(a)

In the ViT architecture, an additional learnable [class] embed is a special marker used to represent the final output of an image classification task. In ViT, this [class] embed is fed into the Transformer Encoder and takes part in the attention calculation along with the other embeddings of the image patch. After being processed by the Transformer Encoder, the final [class] embed contains comprehensive information about the entire image and is used for the final classification task. In the absence of this learnable [class] embed, the embed vectors output by the Transformer Encoder can be aggregated in other ways for classification by using average pooling or maximum pooling of all patches to get a global feature vector and use it for the classification task.

Advantages of using [class] embedding:

Focus on classification tasks: The learnable [class] embedded in the attention mechanism gradually accumulates information about the entire image and is specifically used for classification tasks, avoiding the loss of averaging information from different patches.

More flexible feature aggregation: Compared to average pooling, Transformer Encoder can adaptively provide richer context information for [class] embeddings through the attention mechanism to improve classification accuracy.

Simplified post-processing: The classification output is obtained directly through [class] embedding, and no additional aggregation operations are required for other patch embedding.

4(b)

Positional Embedding is used to provide positional information for the input patch, enabling the Transformer Encoder to understand the relative positions of different patches. In the original NLP Transformer architecture, the input was sequential data, and the order was important, so positional

embedding was added. In ViT, since the image is a two-dimensional structure, the spatial position of the patch is equally critical.

If positional embedding is omitted, the model loses the location information of the patches and cannot distinguish which patches are in which position in the image. This can have the following effects:

Inability to recognize spatial structure: The spatial information of the image (such as shape, edge, position of the object, etc.) is very important for recognizing the object. Without positional embedding, the model cannot effectively distinguish the relative positions of different patches, resulting in degraded classification performance.

Model confusion: Without location information, the model may assume that patches at different locations are interchangeable, which interferes with the classification accuracy of the model.

Therefore, location embedding is essential for ViT models to correctly capture spatial relationships of images in image classification tasks.

4(c)

4(d)

4(e)

i) cls token is a learnable vector with an embedded dimension of $d=16$. Therefore, the number of parameters for the cls token is: 16.

ii) We split the image into 16 patches, each with dimensions of $3 \times 64 \times 64$. After each patch is flattened, it is projected through a linear layer into the embedded dimension $d=16$. **The number of parameters of the Patch Linear Projection is $(3 \times 64 \times 64) \times 16 = 196608$.** Where 16 is the weight parameter of the linear layer, so **16 is also the bias parameter**. Also, each patch has a positional embedding vector with dimension $d = 16$. We have $\text{num_patches} = 16$ patches. Therefore, **the number of parameters of the Position Embedding is: $16 \times 16 = 256$.** In total, **we would have $196608 + 16 + 256 = 196880$ parameters.**

iii) MLP Head is a single-layer fully connected layer that maps the embedded dimension $d = 16$ to the class number C . Let's say we have C categories. The number of parameters for the MLP Head is: $16 \times C + C$. In this case, for $C=8$, we have $16 \times 8 + 8 = 136$. (I actually don't know what this question's background exactly is, if it is still based on 1 that $C=2$, we have $16 \times 2 + 2 = 34$.)

iv)

Each attentional head has a linear projection of Query, Key, and Value that maps the embedded dimension d to query_size . Here $\text{query_size} = d / \text{num_heads} = 16 / 2 = 8$. For each attention head, we have: $(d \times 8 + 8) \times 3 = (16 \times 8 + 8) \times 3 = 408$. There are a total of 2 Attention heads, so **the number of arguments for Multi-Head attention is: $408 \times 2 = 916$.**

The output of multi-head attention is spliced and mapped back to the d dimension through a linear layer. **The number of parameters of this linear layer is: $d \times d + d = 16 \times 16 + 16 = 272$.**

LayerNorm usually has two parameters: weight and bias, each with d parameters. There are two LayerNorms in a Transformer Block, so **the total number of parameters is: $2 \times d \times 2 = 2 \times 16 \times 2 = 64$.**

The MLP consists of two fully connected layers. The first layer maps d to $4d$ and the second layer maps $4d$ back to d . **The number of parameters in the first layer is: $d \times 4d \times 4d = 16 \times 64 \times 64 = 1088$. The number of layer 2 parameters is: $4d \times d \times d = 64 \times 16 \times 16 = 16384$.** Therefore, **the total number of parameters for MLP is: $1088 + 16384 = 17472$.**

In total, we'll have $916 + 272 + 64 + 17472 = 18662$.

v) The entire ViT model contains 2 Transformer blocks, so the total number of parameters in Transformer blocks is: $2 \times 18662 = 37324$.

4(f)

4(g)

The train/validation AUROC performance of the model from the epoch demonstrating the lowest validation loss is epoch 7.

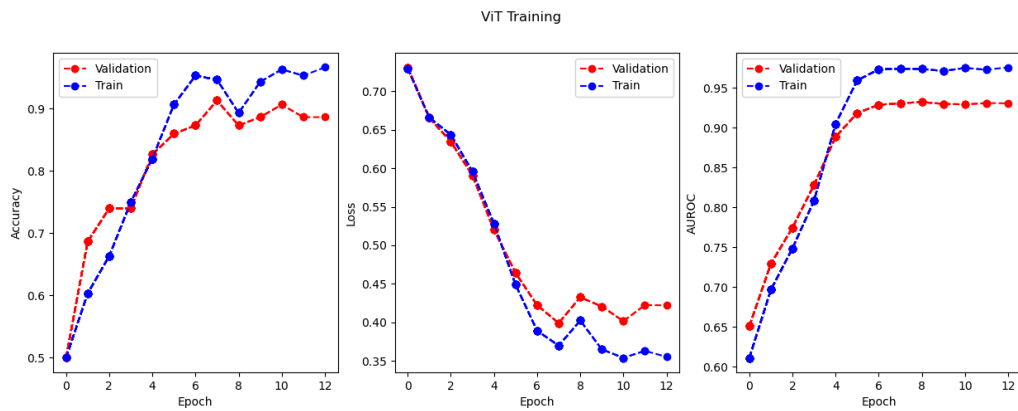
	Training	Validation
AUROC	0.9742	0.9308

```
[(445_p2) mason-wudeMacBook-Pro:probuzhidaogaigaishamingzi:project2 mason]
_wu$ python train_vit.py
loading train...
loading val...
loading test...
Number of float-valued parameters: 7394
Loading ViT...
No saved model parameters found
Epoch 0
    Validation Accuracy:0.5
    Validation Loss:0.7311
    Validation AUROC:0.6508
    Train Accuracy:0.5
    Train Loss:0.7298
    Train AUROC:0.6111
Epoch 1
    Validation Accuracy:0.6867
    Validation Loss:0.6656
    Validation AUROC:0.7292
    Train Accuracy:0.6033
    Train Loss:0.6658
    Train AUROC:0.6966
Epoch 2
    Validation Accuracy:0.74
    Validation Loss:0.6346
    Validation AUROC:0.7746
    Train Accuracy:0.6633
    Train Loss:0.6436
    Train AUROC:0.7486
Epoch 3
    Validation Accuracy:0.74
    Validation Loss:0.5909
    Validation AUROC:0.8283
    Train Accuracy:0.75
    Train Loss:0.5957
    Train AUROC:0.8086
Epoch 4
    Validation Accuracy:0.8267
    Validation Loss:0.5201
    Validation AUROC:0.8885
    Train Accuracy:0.82
    Train Loss:0.5282
    Train AUROC:0.905
Epoch 5
    Validation Accuracy:0.86
    Validation Loss:0.4642
    Validation AUROC:0.9184
    Train Accuracy:0.9067
    Train Loss:0.4489
    Train AUROC:0.9595
Epoch 6
    Validation Accuracy:0.8733
    Validation Loss:0.4224
    Validation AUROC:0.9289
    Train Accuracy:0.9533
    Train Loss:0.3891
    Train AUROC:0.9734
Epoch 7
    Validation Accuracy:0.9133
    Validation Loss:0.399
    Validation AUROC:0.9308
    Train Accuracy:0.9467
    Train Loss:0.3697
    Train AUROC:0.9742
Epoch 8
    Validation Accuracy:0.8733
    Validation Loss:0.4326
```

```

Validation AUROC:0.9323
Train Accuracy:0.8933
Train Loss:0.4024
Train AUROC:0.9738
Epoch 9
Validation Accuracy:0.8867
Validation Loss:0.4206
Validation AUROC:0.93
Train Accuracy:0.9433
Train Loss:0.3648
Train AUROC:0.9713
Epoch 10
Validation Accuracy:0.9067
Validation Loss:0.4015
Validation AUROC:0.9293
Train Accuracy:0.9633
Train Loss:0.3539
Train AUROC:0.9751
Epoch 11
Validation Accuracy:0.8867
Validation Loss:0.4225
Validation AUROC:0.931
Train Accuracy:0.9533
Train Loss:0.363
Train AUROC:0.9728
Epoch 12
Validation Accuracy:0.8867
Validation Loss:0.4221
Validation AUROC:0.9307
Train Accuracy:0.9667
Train Loss:0.3552
Train AUROC:0.9759
Finished Training

```



4(h) & 4(i)

```

[(445_p2) mason-wudeMacBook-Pro:probuzhidaogaigaishamingzi:project2 mason]
_wu$ python test_vit.py
loading train...
loading val...
loading test...
Loading ViT...
Which epoch to load from? Choose in range [0, 12]. Enter 0 to train
from scratch.
>> 7
Loading from checkpoint ./checkpoints/vit/epoch=7.checkpoint.pth.tar
?
=> Successfully restored checkpoint (trained for 7 epochs)
Epoch 7
    Validation Accuracy:0.9133
    Validation Loss:0.399
    Validation AUROC:0.9308
    Train Accuracy:0.9467
    Train Loss:0.3643
    Train AUROC:0.9742
    Test Accuracy:0.59
    Test Loss:0.7232
    Test AUROC:0.5616

```

AUROC performance:

	Training	Validation	Testing
ViT	0.9742	0.9308	0.5616
CNN	0.9963	0.9406	0.6836

Training (training set) : ViT: 0.9742 CNN: 0.9963

Analysis: On the training set, CNN performed slightly better than ViT. The higher training accuracy of CNN may indicate that it is stronger at fitting training data, or it may mean that it is more complex than ViT and suitable for working with this particular data set.

Validation (validation set) : ViT: 0.9308 CNN: 0.9406

Analysis: On the validation set, CNN performs slightly better than ViT, but not by much. This suggests that CNN may have better adaptability to the characteristics of this data set, especially the generalization performance on the validation set is slightly higher than that of ViT.

Testing (test set) : ViT: 0.5616 CNN: 0.6836

Analysis: On the test set, CNN performs significantly better than ViT, and the test accuracy of CNN is about 12% higher than that of ViT. This shows that the generalization performance on the test set is stronger on CNN than on ViT.

This gap may be related to the following reasons:

Training data size and data requirements of Vits: ViT models typically require larger data sets to adequately train because it does not have the convolutional structure of CNN and lacks the ability to extract local features on small data sets.

Advantages of CNN for visual tasks: The CNN model is naturally suitable for processing image data and can efficiently extract spatial local features through convolution kernel. However, ViT mainly relies on self-attention mechanisms to capture long-distance dependencies and may not be fully utilized on smaller data sets.

Feature extraction mechanism of ViT: ViT's self-attention mechanism relies on global information when the amount of data is small or the local features of the image are very small.

5

Here is the deep convolutional neural network I made of classifying Collies and Golden Retrievers. Below is a detailed discussion of the different design decisions I made:

1. Model Architecture

I used a deep convolutional neural network (CNN) with residual connections, with an architecture inspired by ResNet.

Initial Convolutional Layers: The first convolutional layer maps the input RGB image to 64 feature channels with batch normalization and ReLU activation.

Residual Blocks: 3 layers of residual modules are used in the network. Each module contains 2 residual blocks, with a progressively increasing number of channels ($64 \rightarrow 128 \rightarrow 256 \rightarrow 512$), while downsampling is achieved by $\text{stride}=2$. The residual blocks mitigate the gradient vanishing problem in deep networks and ensure effective learning of deep features.

Global average pooling: after the last residual block, the output feature map is reduced to 8×8 using `adaptive_avg_pool2d`. This global pooling approach allows for better adaptation to the input image size.

Fully Connected Layers: Two fully connected layers are used with Dropout for regularization, and the final output is the prediction of the 2 categories.

2. Regularization

In order to prevent the model from overfitting, I use various regularization methods:

Dropout: A 50% dropout is added to the first fully connected layer, which helps reduce overfitting and improves the model's performance on the validation set.

Batch Normalization: Batch normalization is added after each convolutional layer. Batch normalization can speed up the training process, reduce the problem of gradient disappearance, and play a certain role in regularization.

3. Data Enhancement

I apply a variety of data enhancement techniques aimed at increasing the generalization ability of the model, including:

Random cropping and scaling: adding image segments of different sizes and positions through `RandomResizedCrop` to simulate different imaging conditions.

Horizontal Flip: Randomly flips the image horizontally to prevent the model from being biased against image features in a particular orientation.

Color Perturbation: Adjusts the brightness, contrast, saturation, and hue of an image using `ColorJitter` to increase the robustness of the model to changes in lighting.

Random Rotation: Randomly rotate the image to help the model better adapt to different orientations of the target.

4. Hyperparameter Selection

During model training, I adjusted some key hyperparameters to optimize the model performance:

Learning rate: the initial learning rate is set to $1e-4$. Through experiments, I found that this learning rate can converge in a reasonable time while avoiding large oscillations.

Learning Rate Scheduler: The `StepLR` scheduler is used, which reduces the learning rate by 50% every 10 epochs to help the model continue to improve its accuracy in the later stages of convergence.

Batch Size: 32 was chosen as the batch size. Among the different sizes I tried, 32 balanced the memory usage and model training speed.

5. Model Evaluation Criteria

In the model evaluation process, I use the following metrics to judge the strengths and weaknesses of the model:

Accuracy: This is the main evaluation metric of the model, which directly reflects the classification accuracy of the model on the test set.

Loss: I use the loss value of the validation set to evaluate the convergence of the model and use it as a criterion for early stopping.

AUROC (Area Under ROC Curve): AUROC is a better measure of model performance on unbalanced datasets.

With these metrics, I was able to determine which model had better generalization performance.

6. Early Stopping and Patience settings

In order to prevent the model from overfitting during training, I introduced the Early Stopping mechanism:

Patience: Allow up to 5 epochs (patience = 5) before the validation loss decreases, and stop training if this number is exceeded.

Implementation: By calculating the validation loss after each epoch and comparing it with the previous epoch, and increasing the patience counter if there is no significant improvement, triggering an Early Stopping when patience reaches a threshold.

7. Future Improvement Directions

More residual blocks: the network can be further deepened by increasing the number of residual block layers.

Richer data enhancement: Try other data enhancement methods, such as Cutout or Mixup, to further improve the robustness of the model.

Optimize the regularization strategy: In future experiments, Dropout can be introduced in different layers, and even DropBlock can be used in the convolutional layer.

Extra Output:

Which epoch to load from? Choose in range [0, 12]. Enter 0 to train from scratch.

>> 7

Loading from checkpoint ./checkpoints/target/epoch=7.checkpoint.pth.tar?

=> Successfully restored checkpoint (trained for 7 epochs)

Epoch 7

Validation Accuracy:0.8867
Validation Loss:0.3308
Validation AUROC:0.9406
Train Accuracy:0.9667
Train Loss:0.1044
Train AUROC:0.9963

Epoch 8

Validation Accuracy:0.86
Validation Loss:0.4406
Validation AUROC:0.9335
Train Accuracy:0.9267
Train Loss:0.1699
Train AUROC:0.9896

Epoch 9

Validation Accuracy:0.8667
Validation Loss:0.343
Validation AUROC:0.9383
Train Accuracy:0.9533
Train Loss:0.1172
Train AUROC:0.9961

Epoch 10

Validation Accuracy:0.8667
Validation Loss:0.4436
Validation AUROC:0.9346
Train Accuracy:0.9767
Train Loss:0.0873
Train AUROC:0.9976

Epoch 11

Validation Accuracy:0.88
Validation Loss:0.4407
Validation AUROC:0.936
Train Accuracy:0.9867
Train Loss:0.0654
Train AUROC:0.9975

Epoch 12

Validation Accuracy:0.88
Validation Loss:0.3815
Validation AUROC:0.9387
Train Accuracy:0.99
Train Loss:0.0475
Train AUROC:0.9998

Epoch 13

Validation Accuracy:0.8667
Validation Loss:0.4437
Validation AUROC:0.9404
Train Accuracy:0.9733
Train Loss:0.0904
Train AUROC:0.9996

Epoch 14

Validation Accuracy:0.8733
Validation Loss:0.456
Validation AUROC:0.941
Train Accuracy:0.9933
Train Loss:0.0425
Train AUROC:1.0

Finished Training

```

[(445_p2) mason-wudeMacBook-Pro:probuzhidaogaigaishamingzi:project2 mason]
_wu$ python test_cnn.py
loading train...
loading val...
loading test...
Loading cnn...
Which epoch to load from? Choose in range [0, 12]. Enter 0 to train
from scratch.
>> 11
Loading from checkpoint ./checkpoints/target/epoch=11.checkpoint.pth
.tar?
=> Successfully restored checkpoint (trained for 11 epochs)
Epoch 11
    Validation Accuracy:0.8733
    Validation Loss:0.4511
    Validation AUROC:0.936
    Train Accuracy:0.9933
    Train Loss:0.0438
    Train AUROC:0.9996
    Test Accuracy:0.64
    Test Loss:1.4767
    Test AUROC:0.6728

```

	AUROC		
	TRAIN	VAL	TEST
Freeze all CONV layers (Fine-tune FC layer)	0.9724	0.9244	0.7476
Freeze first two CONV layers (Fine-tune last CONV and FC layers)	0.9732	0.9367	0.742
Freeze first CONV layer (Fine-tune last 2 conv. and fc layers)	0.9724	0.9244	0.7476
Freeze no layers (Fine-tune all layers)			
No Pretraining or Transfer Learning (Section 2 performance)	0.9963	0.9406	0.6836